

AP Stats Notes

Franklin Chen

11 November 2024 - idk some time in may

Contents

1	Data Analysis	4
1.1	What is Statistics?	4
1.2	Analyzing Categorical Data	4
1.2.1	One-Variable Categorical Data	4
1.2.2	Two-Variable Categorical Data	5
1.3	Analyzing Quantative Data with Graphs	6
1.4	Describing Quantative Data with Numbers	7
2	Modelling Distributions	9
2.1	Describing Location in a Distribution	9
2.1.1	Transforming Data	9
2.2	Density Curves	10
2.2.1	Normal Distributions	10
2.2.2	Assessing Normality	11
3	Two-Variable Data	12
3.1	Scatterplots and Correlation	12
3.2	Least-Squares Regression	13
3.2.1	Residuals	13
3.2.2	Least-Square Regression Lines	13
3.2.3	Problems with Least-Square Regression Lines	14
3.2.4	Strength of regression lines	14
3.3	Transforming to Achieve Linearity	15
3.3.1	Power Models	15
3.3.2	Transforming with Logarithms	15
4	Collecting Data	16
5	Probability	17
6	Random Variables and Distributions	18
7	Sampling Distributions	19
8	Confidence Intervals	20
9	Significance Tests	21
9.1	Basics of Significance Tests	21
9.1.1	Hypotheses	21
9.1.2	P-values	21

9.1.3	Type I and Type II Errors	23
9.1.4	Power	23
9.1.5	Steps for Significance Tests	23
9.2	Tests About a Population Proportion	24
9.3	Tests About a Difference in Proportions	25
10	Estimating Means	26
10.1	One-Sample t intervals	26
10.2	Two-Sample t intervals for a Difference between Two Means	27
10.3	Paired t interval For a Mean Difference	27
11	Confidence with Means	28
11.1	Tests about a Population Mean	28
12	Inference for Distributions and Relationships	29
12.1	Chi-Square Tests for Goodness of Fit	29
12.2	Inference for Two-Way Tables	31
12.2.1	Chi-square tests for homogeneity	31
12.2.2	Chi-square tests for independence	32
12.3	Inference for Slope	32

using *The Practice of Statistics for the AP Exam: 6th Edition* by Starnes and Tabor

1 Data Analysis

1.1 What is Statistics?

Definition 1 (Statistics). The science of collecting, analyzing, and drawing conclusions from data.

Data is collected from *individuals* about certain *variables*.

Definition 2 (Individual, Variable). **Individuals** are objects described in a dataset. Typically people, but not always.

Variables are attributes that can take different values for different individuals.

For example, *individuals* may be households, and *variables* may be region, number of people, household income, etc. It's important to distinguish between *categorical* and *quantitative* variables:

Definition 3 (Categorical and Quantitative Variables). **Categorical Variables** are variables whose values can be placed into distinct categories.

Quantitative Variables are variables whose values are quantities, typically counts or measurements.

For example, region would be categorical, while household income would be quantitative. *Not all numbers are quantitative*; eg. zip code.

1.2 Analyzing Categorical Data

1.2.1 One-Variable Categorical Data

Definition 4 (Frequency and Relative Frequency Tables). **Frequency Tables** shows the number of individuals that have values of a certain category. **Relative Frequency Tables** shows the proportion or percent of individuals in each category.

Note (relative) frequencies are not data; they summarize data. Bar graphs and Pie Graphs summarize relative frequency tables.

Beware of misleading graphs; we mainly react to the area of each bar, not the actual height.

	Like Skateboards	Do Not Like Skateboards	Totals
Like Snowmobiles	80	25	105
Do not like Snowmobiles	45	10	55
Totals	125	35	160

MathBits.com

Figure 1: An example two-way table with additional summary information.

1.2.2 Two-Variable Categorical Data

Use a two-way table to summarize data about two categorical variables. These tables can be used to answer questions about *marginal, joint, and conditional relative frequencies*.

Margial relative frequencies give the percent or proportion of individuals that have a given value for one categorical variable. For example, the marginal relative frequency of liking skateboards is $\frac{125}{160} \approx 78.125\%$.

Joint relative frequencies give the percent or proportion of individuals that have a specific value for both categorical variables. For example, the joint frequency of liking both skateboards and snowmobiles is $\frac{80}{160} = 50\%$.

Conditional relative frequencies give the percent or proportion of individuals that have a specific value for one categorical variable relative to other individuals with the same other categorical variable. For example, the conditional relative frequency of those who like snowmobiles out of all individuals that like skateboards is $\frac{80}{125} = 64\%$.

These frequencies can be summarized in *side by side bar graphs, segmented bar graphs, or mosaic plots*.

Graphs and these tables can be used to show **association** between two variables. There is association between two variables if knowing the value of one helps to predict the other. For example, knowing that an individual likes skateboards helps predict whether they like snowmobiles ($\frac{80}{125} = 64\%$ vs $\frac{25}{35} \approx 71.4\%$). **ASSOCIATION DOES NOT IMPLY CAUSATION!**




Figure 2: A dotplot showing the distribution of kilometers run by members of the running club.

1.3 Analyzing Quantative Data with Graphs

Dotplots (as shown above) show each individual as a dot above their quantitative data value.

When describing the shape of a dotplot (or other quantitative graphs), *focus on main features*: major peaks, clusters, or gaps. Especially note whether the distribution is roughly symmetric or skewed:

Definition 5 (Symmetric, Skewed). A distribution is roughly **symmetric** if the right side of the graph has roughly the same shape as the left side.

A distribution is **skewed to the right** if the right 'tail' has less values than the left; typically, the left has a peak whereas the right does not. **Left-skewed** definition are defined similarly to right-skewed distributions.

For example, the distribution of the number of kilometers run is right-skewed because the right 'tail' has less values.

Graphs with a single peak are considered *unimodal*, like the dotplot. Distributions with two peaks are considered *bimodal*, and beyond that is considered *multimodal*.

When describing a distribution of quantitative data, use the acronym ROCS: **R**ange (max - min), **O**utliers (clear departures from the data), **C**enter (mean or median), and **S**hape (symmetry, skew, gaps, peaks).

Leaf plots exist. Stem represents first few digits, leaf represents final digit.




Figure 3: An example histogram with a normal distribution.

Histograms are a notable way of displaying quantative data, as they avoid showing individual data points. Histograms divide the variable into many 'bins' (bars), with the height representing the frequency. Smaller bins show more detail at the cost of a less clear pattern.

Don't confuse histograms and bar graphs. Histograms are used for quantative data, while bar graphs are used for qualatative data.

Use percentages when comapring to distributions in order to remove the effect of a larger sample.

1.4 Describing Quantative Data with Numbers

Definition 6 (Mean: \bar{x}, μ). The average of all individual data values. If there are n observations x_1, x_2, \dots, x_n , the sample mean is calculated by

$$\bar{x} = \frac{\sum x_i}{n}$$

The mean of a **sample** is referred to using \bar{x} , while the mean of a **population** is referred to using μ .

Statistics come from **samples** (small subset of population) and **parameters** come from **populations** (all possible samples of what's being tested).

The mean is not **resistant** as it is sensitive to strong outliers in a distribution. The median *is* a resistant measure of the center of the distribution.

Definition 7 (Median). The 'midpoint' of a distribution. Either the middle element (n is odd) or the average of the two middle elements (n is even) in a **SORTED** distribution.

Using both the mean and the median, one can predict the skew of the data. If a distribution is roughly symmetric without outliers, **the mean and median will be similar**. If a distribution is strongly skewed, **the mean will be pulled in the direction of the skew**. (Mean < Median for left-skewed, Mean > Median for right-skewed)

The **range** (max - min) is one way to show the variability of a distribution. Note that the range is *not* resistant.

Definition 8 (Standard Deviation). The **standard deviation** (s_x, σ) measures the 'average' distance of the values in a distribution from the mean. Standard deviation is calculated by

$$s_x = \sqrt{\frac{\sum (x_i - \bar{x})^2}{n - 1}}$$

The squared stdev is known as **variance** (s_x^2, σ^2). Remember, s_x refers to a sample while σ refers to a population. Larger stdev indicates greater variation, but is not a resistant measure of variability. **Stdev measures variance around the mean; if the mean is skewed, so will stdev!**

The **Interquartile Range (IQR)** is another way to measure variance, using $IQR = Q_3 - Q_1$ where Q represents the quartiles. IQR can be thought of as the range of the 'middle half' of the distribution. *IQR is a resistant measure.*

Lower Outliers < $Q_1 - 1.5 \times IQR$ or High Outliers > $Q_3 + 1.5 \times IQR$

boxplots and the five-number summary = min, Q1, median, Q3, and max exist. **Boxplots don't show gaps, clusters, or multiple peaks.**

be careful with language- 'skews' is a shape, IQR and range are single numbers (no 'in the middle of the IQR')

2 Modelling Distributions

2.1 Describing Location in a Distribution

Definition 9 (Percentile). The p th percentile of a distribution is the value with $p\%$ of observations less than (*or equal, depending on who you ask*) than it. Note that this distinction of whether or not to include "or equal" only matters for discrete variables.

For example, in a class of 25 students, if a student gets a score greater than or equal to 21 other students, then they would be at the 84th percentile in the class's test score distribution ($\frac{21}{25} \approx 84\%$). **An observation is never "in" a percentile; an observation is AT a percentile** (percentile is just a number; like IQR and range). Also, definition, $Q_1 \approx 25$ th percentile, Q_2 (median) ≈ 50 th percentile, and $Q_3 \approx 75$ th percentile. Percentiles can be graphically shown in a **cumulative relative frequency graph**, where the y-coordinates of points are graphed based on their percentile.

Definition 10 (Standardized Score (z-score)). How many standard deviations an individual value is from the mean. For a value p , mean μ , and stdev σ :

$$z = \frac{p - \mu}{\sigma}$$

z-scores provide a way to objectively compare measurements from two distributions while still considering means and variability.

2.1.1 Transforming Data

Effect of adding a constant a :

- Adds a to measures of center and location (mean, median, quartiles, min, max)
- Does not change measures of variability (range, IQR, stdev)
- Does not change the shape of the distribution (percentile; translation along axis)

Effect of multiplying by constant b :

- Multiplies measures of center and location by b (mean, median, quartiles, min, max)
- Multiplies measures of variability by $|b|$ (range, IQR, stdev)

- Does not change the *overall* shape of the distribution (percentile; like squish or squash)

z-scores transform any distribution into one with mean 0 and stdev 1, but with the same original shape.

2.2 Density Curves

Definition 11 (Density Curves). A curve that models the distribution of a quantative variable with a curve that

- Is always on (or above) the horizontal axis
- Has an area of exactly 1 underneath

The area under the curve within any interval of values estimates the proportion of all observations that fall into this interval.

No set of quantative data is fully described by a density curve. Density curves are approximations that are easy to use by smoothing out small irregularities in the data. Similarly to finite distributions, distributions will have shapes, often with skews and peaks.

The median of the density curve splits the curve into two equal-area halves of area = 0.5. The mean of the density curve is harder to define. For a density curve described by a function $p(x)$, the mean is given as the below. (Intuitively, this is the idea of a 'weighted average'; weight being its relative density- $p(x)$, in this case- being extended to a continous distribution.)

$$\text{mean of } p(x) = \int_{-\infty}^{\infty} xp(x)dx$$

While the mean-median location principles still apply, they aren't too useful as it's relatively hard to locate the mean of a curve by simply looking at it.

2.2.1 Normal Distributions

As $\lim_{n \rightarrow \infty} n$, binomial distributions will approach a **Normal distribution** (see Chapter 6 or something), which are described using Normal curves.

Definition 12 (Normal Curve). Any normal distribution is described by a symmetric, single-peaked, bell-shaped density curve. Its center is equal to the mean and the median, and the stdev represents the 'width' of the curve. Any normal distribution is fully described by its mean and stdev. A normal

curve's mean is its peak, while the stdev are its inflection points (symmetric around the mean).

Normal distributions with $\mu = 0$ and $\sigma = 1$ is known as a **standard Normal distribution** (which is the same as any other Normal distribution normalized into z-scores). *Remember that mean and stdev only fully describe Normal distributions!*

All normal distributions follow the **empirical rule**: 68%, 95%, and 99.7% of all observations will fall within 1, 2, and 3 stdev around the mean respectively. `normalcdf(upper, lower, mean, stdev)` can be used to find the area under the normal curve (with a given mean and stdev) between upper and lower. Similarly, `invNorm(area to the left, mean, stdev)` calculates the area's percentile value given the mean and stdev. **When using calculator functions like normalcdf, make sure to 1) label what the inputs mean and 2) answer the question asked AS A SENTENCE!**

2.2.2 Assessing Normality

uhh i'll do it later

3 Two-Variable Data

3.1 Scatterplots and Correlation

When analyzing two sets of data, we usually hope to establish a cause-effect relationship between the **independent / explanatory variable** and the **dependent / response variable**. We want to show that knowing the value of the independent value will help predict the dependent variable's value. **Just because knowing the independent value helps us predict the dependent value does not imply direct causation; there may be some other variable influencing both!**

We can analyze two-variable using a scatterplot. To describe a scatterplot, consider the following:

- **DIRECTION:** Does the dependent variable increase (*positive association*) or decrease (*negative association*) with the independent variable? Or does it look like there's no relation between the two? (*no association*)
- **FORM:** Is the data linear or nonlinear? Nonlinear data still shows a pattern, although it's harder to describe.
- **STRENGTH:** How closely does data follow the predicted pattern? Not very much (*strong association*) or a lot (*weak association*)? Or somewhere between? (*moderate association*)
- **OUTLIERS:** Look at distinct departures from the described pattern and distinct clusters of points.

We can directly quantify correlation for linear relationships using the **(linear) correlation r** .

Definition 13 (Correlation r). A measure of the direction and strength of a linear association. Always between -1 and 1, with values closer to each extreme indicating stronger relationships ($r = \pm 1 \rightarrow$ perfectly linear relationship, $r = 0 \rightarrow$ no linear relationship). Positive r suggests positive association, while negative r suggests negative association.

$$r = \frac{1}{n-1} \sum \left(\frac{x_i - \bar{x}}{s_x} \right) \left(\frac{y_i - \bar{y}}{s_y} \right) = \frac{1}{n-1} \sum z_x z_y$$

On a surface level, the formula for r standardizes each variable to avoid the issue of different linear slopes. Other than that, don't ask why this formula works; it's wayyyy too much of a hassle to reasonably prove.

Obviously, r only works if both variables are quantitative; use the word *association* if either variables are categorical. r doesn't differentiate between explanatory and response variables; r will be the same even if you swap the two sets. r also has no units; it's just a scalar.

The linear correlation should only be used when talking about LINEAR relationships: a perfectly quadratic relationship will not give $r = 1$. Linear correlation does not care about form; even if something is curved, linear correlation only measures its 'closeness' to a linear line. **Linear correlation is not resistant; it is GREATLY influenced by extreme points.**

3.2 Least-Squares Regression

Definition 14 (Regression Line). A line of form $\hat{y} = a + bx$ which models how the dependent variable y changes as the independent x changes. (\hat{y} = predicted value of y for given x) **Regression lines require that we consider the x as the independent variable.**

Alternatively, we can think of \hat{y} as the average price for a sample of individuals with the same x value (assuming that our pattern is true). Beware of **extrapolation**: using a regression line to predict outside of the interval of x -values used to obtain the line. The further the values are from the data used to generate the line, the less reliable the prediction is. *Just because regression lines can predict extreme values doesn't mean those values will ever happen!*

3.2.1 Residuals

To determine if a regression line is a good prediction, plot the **residuals** ($y - \hat{y}$ for all data points in the set). No pattern suggests that the regression model is accurate, while a leftover non-zero pattern suggests that this model isn't appropriate. (Close to the x -axis = strength)

Intuitively, we look at the remaining pattern because **form of residual plot = form of association - form of model**. As such, no residual plot form suggests that form of association = form of model (good), while a residual plot form suggests that form of association \neq form of model (bad!!!).

3.2.2 Least-Square Regression Lines

A good regression line minimizes the **sum of the squared residuals**. (Why not the residuals themselves? The sum of all the residuals = 0, so we square it to avoid sign- similar to stdev.) The line that minimizes the squared residuals is known as the **least squares regression line**.

The least-square regression line $\hat{y} = a + bx$ can be calculated as follows:

$$b = r \frac{s_y}{s_x}$$

$$a = \bar{y} - b\bar{x}$$

The formula for the slope can be thought of as "normalizing" the correlation to fit with the specific scenario using the coefficient $\frac{s_y}{s_x}$. (as far as i can tell there's no super intuitive way to think about this. sorry)

The formula for the y-intercept comes from the fact the least-squares regression line will always pass through (\bar{x}, \bar{y}) . Using simple algebra we rearrange for a.

This is called *regression to the mean* because something something predicted values "regress" to the mean somehow.

3.2.3 Problems with Least-Square Regression Lines

Similar to correlation, least-square regression lines are 1) only used to describe linear relationships, 2) not resistant, and 3) only describe patterns in the data, not direct causation.

2 suggests that points with relatively high or low x-values have **high leverage** over the regression line versus other points in the data set. These **outliers** (original def + high residual) are **influential points**; if removed, they substantially change the regression line or measures of strength (see 3.2.4).

3.2.4 Strength of regression lines

To determine the strength of a regression line, we have two methods: the **standard deviation of the residuals** or the **coefficient of determination** r^2 .

The standard deviation of the residuals (s) is pretty obvious: calculate the stdev of the data relative to the prediction. (why $n - 2$? idk it looks rllly complicated)

$$s = \sqrt{\frac{\sum (y_i - \hat{y}_i)^2}{n - 2}}$$

On the other hand, the coefficient of determination measures the percent reduction in the sum of squared residuals when using the regression line vs. using $\hat{y} = \bar{y}$. In other words, it represents the percent of variability in the response variable that is accounted for by the least squares regression line. ($1 - r^2$ can be thought of as the % of variability that's caused by other factors)

than the independent variable.) (It's also equal to the square of the linear correlation, hence the symbol r^2 . why? idk)

$$r^2 = \frac{\sum(y_i - \bar{y})^2 - \sum(y_i - \hat{y})^2}{\sum(y_i - \bar{y})^2}$$

3.3 Transforming to Achieve Linearity

The techniques discussed in 3.1 and 3.2 (obviously) only apply to purely linear relationships. In some special cases, we may be able to **transform the data** from a non-linear pattern into a linear pattern to apply our techniques.

3.3.1 Power Models

Definition 15 (Power Model). Any regression model that takes the form $\hat{y} = ax^p$. Scenarios modeled by these regression models should be expected to do so by geometry (eg. area proportional to square, volume proportional to cube) or some other factor like physics.

While power models (by definition) describe a non-linear relationship between x and y , there will always be a linear relationship between x^p and y . As such, if we raise x to the p th power or take the p th root of y there should be a linear relationship between x^p and y (or x and $\sqrt[p]{y}$).

However, this all assumes p is known. If p isn't known, you could either guess and check (cringe) or use logarithms.

3.3.2 Transforming with Logarithms

If a scenario is modeled by $y = ax^p$, then $\ln(y) = \ln(ax^p) \rightarrow \ln(y) = \ln(a) + p\ln(x)$. Thus, there is a linear relationship between $\ln(x)$ and $\ln(y)$. (*Note that the power p now becomes the slope of the line!*) If the resulting graph turns out to be linear, we can fit a linear regression line to the graph and then use the relation $\hat{y} = e^{a+b\ln(x)}$ to predict values. (This relationship is the same as the power model: $\hat{y} = e^a \times e^{b\ln(x)} = e^a x^b$; e^a is a constant.)

If a scenario is modeled by $y = ab^x$ (an **exponential model**), then $\ln(y) = \ln(ab^x) \rightarrow \ln(y) = \ln(a) + \ln(b)x$. Thus, there is a linear relationship between x and $\ln(y)$. If the resulting graph turns out to be linear, we can fit a linear regression line to the graph and use the relation $\hat{y} = e^{a+bx}$. (This relationship is the same as the exponential model: $\hat{y} = e^a \times e^{bx} = e^a \times (e^b)^x$; e^a and e^b are constants.)

If unsure if an exponential or power model is better, consider their residual plots. If both models have similarly random residual plots, choose the model with the largest r^2 .

4 Collecting Data

5 Probability

6 Random Variables and Distributions

7 Sampling Distributions

8 Confidence Intervals

9 Significance Tests

Significance Tests are like the opposite of confidence intervals. Instead of using a statistic to find a parameter, significance tests use statistics to test claims about a parameter.

Definition 16 (Significance Test). A formal procedure for using observed statistics in order to decide between two competing claims (*hypotheses*) about parameters.

9.1 Basics of Significance Tests

9.1.1 Hypotheses

Definition 17 (Null Hypothesis, H_0). A claim about a parameter that we weigh evidence **against** in a significance test. Usually a statement of 'no difference' (as claimed).

Definition 18 (Alternative Hypothesis, H_a). The claim that we are trying to find evidence for. Directly contradicts the null hypothesis.

This "Null Hypothesis" and "Alternative Hypothesis" can be thought as trying to prove someone "guilty" or "not guilty."

For example, if a player claims they're a 80% free throw player, the null hypothesis would be $H_0 : p = 0.80$, and the alternative hypothesis would be $H_a : p < 0.80$.

The alternative hypothesis is **one-sided** because we suspect that the player makes less than 80% of his free throws. If we believe that it's equally plausible that they make more than 80% of their free throws, then we would use a **two-sided hypothesis**- $H_a : p \neq 0.80$.

Hypotheses express our beliefs before looking at the data. Modeling a hypothesis around data shows nothing. (11.1) **Just because a test is statistically significant does NOT mean that it is PRACTICALLY significant.** Focus on the P-value and better yet, the data itself. Graph your data and see if differences or outliers are significant, or if there's another pattern emerging. Generally, for practical importance, it's much better to give a confidence interval versus a significance test.

9.1.2 P-values

Definition 19 (P-value). The probability of getting the values observed in the data under the assumption that the null hypothesis H_0 is true.

Small P-values provide convincing evidence for the alternative hypothesis because small values suggest the observed result is unlikely to happen when the null hypothesis is true. Similarly, large P-values provide convincing evidence for the null hypothesis because large values suggest the observed result is likely to happen due to chance if the null hypothesis is true.

In terms of probability notation: P-value = $P(\text{observed data} \mid \text{null hypothesis is true})$.

For two-sided tests, we look at the distance between the null hypothesis and the observed data. For example, if $H_0 : p = 0.5$, $H_a : p \neq 0.5$ and $\hat{p} = 0.65$ (observed proportion), then the P-value = $P(\hat{p} \leq 0.35 \text{ or } \hat{p} \geq 0.6 \mid p = 0.5)$. We look at $\hat{p} \leq 0.35$ because $|p - 0.35| = |p - 0.65|$.

Based on the P-value, we make a conclusion about data:

- If the P-value is small (unlikely to happen by chance), then we "reject H_0 " and conclude that there is convincing evidence for H_a (in context).
- If the P-value is large (likely to happen by chance), then we "fail to reject H_0 " and conclude that there is not convincing evidence for H_a (in context).

How small does a P-value have to be in order to reject H_0 ? We use a given **significance value** for this boundary.

Definition 20 (Significance Level, α). The value that we use as a boundary for deciding whether a P-value is significant enough to disqualify the null hypothesis. α should be stated before data is produced (*cherry-picking*).

$\text{P-value} < \alpha \Rightarrow \text{reject } H_0 \Rightarrow \text{convincing evidence for } H_a \text{ in context}$

$\text{P-value} > \alpha \Rightarrow \text{fail to reject } H_0 \Rightarrow \text{not convincing evidence for } H_a \text{ in context}$

If P is less than the significance level, we say that the result is "statistically significant at the $\alpha = \text{----}$ level." Alternatively, "the results were significant ($P = 0.03 < \alpha = 0.05$)."
Keep in mind a P-value is more informative than a statement of significance!

NEVER "accept H_0 " or conclude that H_0 is true! Always use 'reject' or 'fail to reject'!

(11.1) **Beware of multiple analyses.** A significance test of $\alpha = 0.05$ is reasonable enough to prove something on its own. But if you take 20 tests of $\alpha = 0.05$ and record if any of them show your conclusion- well, that doesn't really show anything! "Searching for patterns is fine. But performing every conceivable significance test on a data set until you obtain a statistically significant result is not." (This is known as *P-hacking* or *data dredging*.)

9.1.3 Type I and Type II Errors

When drawing a conclusion from a significance test, our conclusion may be wrong. There are two types we can make with the conclusion process, helpfully named "Type I" and "Type II" errors. Only one type of error is possible at once.

Definition 21 (Type I and II errors). A **Type I error** occurs when we reject H_0 when H_0 is true; the data gives convincing evidence that H_a is true despite being false.

A **Type II error** occurs when we fail to reject H_0 when H_a is true; the data fails to give convincing evidence that H_a is true despite it being true.

Note $P(\text{Type I error}) = \alpha$. However, significance is inversely proportional to Type II error: as significance decreases, $P(\text{Type II error})$ increases.

9.1.4 Power

Definition 22 (Power). The **power** of a test is the probability that the test finds convincing evidence for H_a given that the parameter is a specific value that does follow H_a . In other words:

$$\text{power} = P(\text{reject } H_0 \mid H_0 \text{ is false})$$

Power depends on a specific value. For example, if the power of a test to detect $p = 0.08$ is 0.29 given $H_0 : p = 0.10$, if the true proportion in the population is $p = 0.08$, there is a 0.29 probability that researchers will find convincing evidence for H_a .

Note $\text{Power} = 1 - P(\text{Type II Error})$, or written alternatively, $P(\text{Type II Error}) = 1 - \text{Power}$. Also note that the power of a test increases with higher sample number, higher significance level, further null and alternative parameter values, and 'wise choices when collecting data' (reducing variability). *Larger significance values reduce the chance for Type II error, but increase the chance for Type I error.*

9.1.5 Steps for Significance Tests

Overall, significance tests can be summarized in four steps:

1. **STATE** the hypotheses, significance level, and parameters.
2. **PLAN** the appropriate inference method and check conditions.
3. **DO** the significance test (see Section 7.2):

- Give the sample statistic(s)
 - Calculate the standardized test statistic
 - Find the P-value.
4. **CONCLUDE** whether or not you believe in the null or alternate hypotheses, with reference to P-values, in the context of the problem. Make sure to reference the parameter, not the sample statistic!

9.2 Tests About a Population Proportion

Like confidence intervals, significance tests should satisfy several conditions; random sampling, no bias, 10% if applicable ($n < 0.10N$), and large counts ($np_0 \geq 10$ and $n(1 - p_0) \geq 10$). Note that the large counts condition uses p_0 : the number of successes and failures *assuming the null hypothesis is true* are both greater than or equal to 10.

To conduct a **one-sample z test for a proportion** (CITE THIS WHEN FOLLOWING THE SIGNIFIANCE TEST STEPS) about a proportion, look at the normal distribution *assuming the null hypothesis is true*.

From Section 7.2 (add ref later), $\mu_{\hat{p}} = p_0$ and $\sigma_{\hat{p}} = \sqrt{\frac{p_0(1-p_0)}{n}}$.

Using this, we standardize the statistic with respect to the null value to get the **standardized test statistic**; how many stdev units away the sample statistic is from what we would expect, assuming the null hypothesis is true. **YOU HAVE TO PUT THE STANDARDIZED TEST STATISTIC ON THE TEST!**

$$\text{standardized test statistic} = z_{sts} = \frac{\hat{p} - \mu_{\hat{p}}}{\sigma_{\hat{p}}}$$

Using this, we can directly find the P-value using `normalcdf` with the appropriate values:

- $P(z > z_{sts})$ if $H_a : p > p_0$
- $P(z < z_{sts})$ if $H_a : p < p_0$
- $P(z > |z_{sts}|) + P(z < -|z_{sts}|)$ if $H_a : p \neq p_0$

Keep in mind that just because the P-value is low *does not prove that H_0 is false*; sample proportions may be small due to sampling variability, or we have made a Type I error somehow. **AVOID JUMPING TO CONCLUSIONS**: just because a sample statistic looks unlikely, doesn't mean its P-value is!

It should be noted there is a link between *two-sided tests* and confidence intervals: a K% confidence interval gives an approximate set of p_0 's that would not be rejected by a two-sided test at the $\alpha = \frac{K}{100}$ level, with all other values being rejected.

Intuitively, if a confidence interval for p using *the same* \hat{p} does not include p_0 , as the distribution (test) about p_0 has a similar stdev (due to the same α) the distribution is unlikely to include the actual value of p .

9.3 Tests About a Difference in Proportions

Tests about difference in proportions work very similarly to tests about single proportions. For null hypothesis $H_0 : p_1 - p_2 = 0$, large counts can't directly be applied; we need estimate the true difference in proportions p using a weighted average of the two proportions:

$$\hat{p}_C = \frac{\text{number of successes in both samples combined}}{\text{number of individuals in both samples combined}} = \frac{n_1\hat{p}_1 + n_2\hat{p}_2}{n_1 + n_2}$$

We use this pooled proportion (*even for non-zero differences*, although the given formulas have to be changed) because the combined proportion provides *the best estimation of the true difference in proportions*. In the one-sample z test for a proportion, we used p_0 because we're looking at the distribution relative to it. For **two-sample z tests for the difference between two proportions**, we look at the distribution 'relative' to the true proportion; which our best guess is \hat{p}_C . (something something using each individual \hat{p}_i causes too much error + causes some error in the underlying math when we're testing for 0 diff)

In this case, we use the combined proportion for large counts *for both samples*: that is, $n_1\hat{p}_C, n_1(1 - \hat{p}_C), n_2\hat{p}_C, n_2(1 - \hat{p}_C)$ all need to be greater or equal than 10. Using the combined proportion works for a non-zero difference in the null hypothesis, although is not strictly necessary (large counts for each individual distribution can be used).

The standardized test statistic also uses the combined estimate (for both proportions are equal; otherwise, each individual statistic can be used):

$$z = \frac{(\hat{p}_1 - \hat{p}_2) - 0}{\sqrt{\frac{\hat{p}_C(1-\hat{p}_C)}{n_1} + \frac{\hat{p}_C(1-\hat{p}_C)}{n_2}}}$$

10 Estimating Means

10.1 One-Sample t intervals

To create a confidence interval for a mean, we can use the same point estimate \pm margin of error = (critical value)(SE of statistic) idea. However, this is slightly flipped, as (from Section 7) the stdev of a sample distribution relies on the stdev of the population itself. If we know the stdev of the population, we can the familiar *one-sample z interval for a population interval*:

$$\hat{\mu} = \bar{x} \pm z^* \frac{\sigma}{\sqrt{n}}$$

Using the sample stdev alone is not a good enough approximation. For an example, using the sample stdev to estimate a 99% confidence interval would result in a $\approx 92.5\%$ confidence interval. This is because there is significant variation in the sample stdev itself.

To account for this, we use the larger *t-distribution* instead. The t -distribution is like the z -distribution (measures how many stdevs needed for desired confidence), but has higher area in the tails to compensate for the errors in the sample stdev. Notably, t -distributions vary depending on the sample size; that is, they have a specific *degree of freedom df* . (Use `invT` to calculate the specific t^* value needed.) When using t -distributions, make sure to set $df = n - 1$. Thus, for the *one-sample t interval for a population mean*,

$$\hat{\mu} = \bar{x} \pm t^* \frac{s_x}{\sqrt{n}}$$

The random (and 10%) clauses still apply, although the Large Counts condition is slightly modified due to the central limit theorem:

- If the sample size is small ($n < 30$), then graph the sample data (dotplot is preferred, although boxplot works it's kinda funky) and consider whether the data plausibly comes from a normal population. No strong outliers or skew indicates yes. *You will have to show your graph when assesing the normality of the data.*
- If the sample size is big ($n \geq 30$), then you can still use a one-sample t interval for a population mean (although the true confidence level will be lower than presented).

The `TInterval` function on the calculator can be used to simulate a one-sample t interval for a population mean if the above are satisfied.

As $df \rightarrow \infty$, the t -distribution approaches the normal distribution. Intuitively, as sample size increases, the sample stdev will get closer to the population stdev. The quantity $\frac{s_{\bar{x}}}{\sqrt{n}}$ is known as the *standard error of the sample mean* ($s_{\bar{x}}$).

10.2 Two-Sample t intervals for a Difference between Two Means

Like two-sample z intervals, these require the one-sample requirements to be fulfilled for the populations, and follow the same idea (sqrt the variance):

$$\hat{\mu}_{x_1-x_2} = (\bar{x}_1 - \bar{x}_2) \pm t^* \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}$$

However, the degrees of freedom are significantly more confusing to decide. There are two practical options to decide the df to use:

- **Technology:** Use the formula below to calculate df . Note that in this case, df is usually not a whole number. Results in confidence intervals that are approximately the stated confidence levels.

$$df = \frac{\left(\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}\right)^2}{\frac{1}{n_1-1}\left(\frac{s_1^2}{n_1}\right)^2 + \frac{1}{n_2-1}\left(\frac{s_2^2}{n_2}\right)^2}$$

- **Conservative:** $df = \min(n_1 - 1, n_2 - 1)$. Note this method causes a margin of error as large if not larger than is needed for the desired confidence level.

10.3 Paired t interval For a Mean Difference

(also known as a one-sample t interval for a mean difference)

For comparing data which interviews the "same" (or highly related) individuals (known as **paired data**), "combine" the data into one dataset by subtracting each individuals' values, and analyze that using the tactics from section 10.1. (n = the number of pairs)

The requirements are the same as the 1-sample t interval, although the meaning of 10% ($n_{diff} < 0.10N_{diff}$) is different: we suspect that the data is not independent, so 10% guarantees that the differences are independent.

11 Confidence with Means

tl;dr the same thing as Section 9 but with means this time (it works how you think it does)

11.1 Tests about a Population Mean

Null hypotheses, alternative hypotheses, and significance levels exist and mean the same thing as they did in Section 9. Yipee!

While the random and 10% conditions still apply, the large counts condition is modified to account for the Central Limit Theorem, like described in Section 10.1. To reiterate; if the $n < 30$, check whether the data is approximately Normal (no major skew or outliers) before performing tests; otherwise, everything proceeds as normal (haha). Dotplots are recommended to check for Normality due to ease, but other graphs work too.

As before, we need to calculate a standardized test statistic and a P-value.

$$\text{standardized test statistic} = z = \frac{\bar{x} - \mu_0}{\sigma_{\bar{x}}} = \frac{\bar{x} - \mu_0}{\frac{\sigma}{\sqrt{n}}}$$

Obviously, σ for the population is rarely known. So we use the *standard error of \bar{x}* once again:

$$t = \frac{\bar{x} - \mu_0}{\frac{s_x}{\sqrt{n}}}$$

Instead of using the Normal distribution (like you would when σ is known), you will need to use a *t-distribution*, like Section 10 describes. Same **tcdf**, $df = n - 1$, etc. Remember, if you're lazy: $P(t \geq k) = P(t \leq -k)$ for both Normal and t distributions. No calculator for t-distribution? Use the tables, but be wary to give a range between the floor and ceil of the df.

The same link between two-sided confidence intervals and confidence intervals exist. Interestingly, there's also a relation between one-sided tests and confidence intervals for a population mean.- a one sided test is the same as a confidence interval with $C = 1 - 2\alpha$.

Difference in means and paired means work exactly like you think they do. Remember to use \bar{x}_{diff} and s_{diff} for paired data. Remember to use $df = \min(n_1 - 1, n_2 - 1)$ or technology; don't use the sum.

(idk put this somewhere else) **Use one-sample t procedures when testing data collected in pairs. Use two-sample difference t procedures when testing data collected from two populations independently.**

12 Inference for Distributions and Relationships

12.1 Chi-Square Tests for Goodness of Fit

Obviously, the significance tests and confidence intervals discussed in sections 7 - 11 only apply to single continuous variables. This obviously raises the question: *what if you want to consider multiple variables at once?*

Say, for example, several mutually exclusive proportions, like dice rolls. Of course, you could run a single-variable test on every single sample proportion to see if it matches up with what you expect.

Besides the obvious problem of being fairly inefficient, this also raises the problem of multiple tests indirectly leading to **P-hacking** (running several tests until you get a statistically significant result). Additionally, it doesn't tell you how likely it is to get a random sample with the given sample distribution assuming the claimed proportions are true. Instead, we need to consider all proportions at once, using a *chi-square test for goodness of fit*.

For our null hypothesis, we assume that the claimed distribution is true. This leads to the obvious alternate hypothesis: the true distribution is not the same the claimed distribution. For example, dice rolls:

- H_0 : The distribution of dice rolls is the same as the claimed distribution. In other words: $p_i = \frac{1}{6}$ where p_i represents the proportion of dice rolls that land i th side up.
- H_a : The distribution of dice rolls is not the same as the claimed distribution. At least two of these proportions differ than the values stated in the null hypothesis (2 cause $\sum p_i = 1$)

Like other significance tests, we compare our *expected values* (assuming the null hypothesis is true) to our experimental data. Say we roll a dice 60 times. Then, our expected count for categorical variable i in the distribution is $np_i = (60)(\frac{1}{6}) = 10$ rolls to come up for each side of the die. (*np_i may not be an integer, and shouldn't be rounded to one.*)

To quantify how far our experimental data is compared to our expected data, we calculate the *chi-square test statistic*.

Definition 23 (Chi-square test statistic, χ^2). A measure of how far each experimental count differs from the expected count for each categorical variable. (Summed over all possible values for the categorical variable, k .)

$$\chi^2 = \sum \frac{(\text{Observed Count} - \text{Expected Count})^2}{\text{Expected Count}} = \sum_{i=0}^k \frac{(n\hat{p}_i - np_i)^2}{np_i}$$




Figure 4: A probability density plot showing chi-square distributions with $df = 1, 2, 3, 4, 5$.

Make sure to use the COUNTS, and not the PROPORTIONS when calculating χ^2 . We square the difference for similar reasons to σ : to get rid of sign (i think). We divide by the expected count to encode its relative significance, instead of prioritizing extremely high differences.

Similar to previous significance tests, we look at the *chi-square distribution* to generate a P-value. Chi-square distributions are right-skewed non-negative distributions (due to the square) that also take in degrees of freedom (df) = number of categorical variables - 1. The mean of a chi-squared distribution is equal to df and has a singular peak at $df - 2$ ($df > 2$).

Similar to previous significance tests, the random and 10% conditions still apply. Large counts is slightly modified; all *expected* counts should at least be 5.

Use the chi-squared cdf function on your calculator to generate a P-value. (Remember, your P-value is the area to the RIGHT of your χ^2 value; your lower value should be your test statistic, not the upper!) From there, use the stated significance level to either reject or fail to reject the null hypothesis. Remember that just because we fail to reject the null hypothesis does NOT mean that it is true. We simply don't have data to reject the null hypothesis.

If you are asked how the distribution is different than the predicted distribution, note which categories that contribute the most to χ^2 . Then, describe how the observed and predicted counts for these categories differ, with reference to direction. (less or more?)

Gender	Union Member	Represented by Union	Non-Union	Total
Male	76	84	527	687
Female	67	75	493	635
Total	143	159	1,020	1,322

Figure 5: A two-way table comparing gender and union participation.

12.2 Inference for Two-Way Tables

Two-sample z tests (discussed in section 9.3) allow us to compare proportions in two populations or two treatments. *Chi-square tests for homogeneity* allow us to compare distributions of a single categorical variable across several populations or treatments. *Chi-square tests for independence* allow us to analyze whether there's evidence of an association between two variables.

12.2.1 Chi-square tests for homogeneity

To begin, create a two-way table like shown in Figure 5. Generally, our null hypothesis says that there is *no difference* in the true distribution of a categorical variable in the populations of interest (or treatments in an experiment). In this case, our null hypothesis would be H_0 : There is no difference in the true distributions of union participation if the worker is male or female. Similarly, the alternate hypothesis says that *there is a difference, but does not specify what that difference is*. (H_a works for ANY difference in distribution.) In this case, H_a : There is a difference in the true distributions of union participation if the worker is male or female.

(doesn't really belong here) **Performing multiple tests on the same data increases the probability of making a Type I error in at least one test. This is P-hacking.**

This test has similar requirements to that of chi-squares for goodness of fit: random, 10%, and large counts- all expected values are greater than five. How expected values are calculated is slightly different, as we need to account for different sample sizes. For example, from the table, 687 out of 1322 workers are male. If gender has no impact on union participation, then the proportion of each category of union participation should be the same $\frac{687}{1322} \approx 0.520$. Because 143 workers are union members, we would expect $143 \times \frac{687}{1322} \approx 74.31$ union members to be male, on average. Generally, the expected value for any given cell is equal to $\frac{\text{row total} \times \text{col total}}{\text{table total}}$.

	Union Member	Represented by Union	Non-Union	Total
Male	74.312	82.627	530.061	687
Female	68.688	76.373	489.939	635
Total	143	159	1020	1322

Figure 6: The compliment to Figure 5, showing the expected value for each cell. why doesn't matplotlib allow me to export a higher quality image istfg

Note that in the expected value table, the row and column sums are still the same, by design. From here, you can calculate the test statistic χ^2 and P-value (using $df = (\text{number of rows} - 1) * (\text{number of cols} - 1)$) as discussed in section 12.1. **Remmber to cite the name of the test (χ^2 test for homogeneity),the test statistic, df, and P-value when conducting the test!**

12.2.2 Chi-square tests for independence

While chi-square tests for homogeneity test *if the distributions of two different populations are the same*, chi-square tests for independence test *if knowing one variable helps predict the value of the other*.

12.3 Inference for Slope