

AP Stats Notes

Franklin Chen

11 November 2024 - idk some time in may

Contents

1	Data Analysis	3
1.1	What is Statistics?	3
1.2	Analyzing Categorical Data	3
1.2.1	One-Variable Categorical Data	3
1.2.2	Two-Variable Categorical Data	4
1.3	Analyzing Quantative Data with Graphs	5
1.4	Describing Quantative Data with Numbers	6
2	Modelling Distributions	8
3	Two-Variable Data	9
4	Collecting Data	10
5	Probability	11
6	Random Variables and Distributions	12
7	Sampling Distributions	13
8	Confidence Intervals	14
9	Significance Tests	15
9.1	Basics of Significance Tests	15
9.1.1	Hypotheses	15
9.1.2	P-values	15
9.1.3	Type I and Type II Errors	16
9.1.4	Power	17
9.1.5	Steps for Significance Tests	17
9.2	Tests About a Population Proportion	18
9.3	Tests About a Difference in Proportions	18
10	Estimating Means	20
11	Confidence with Means	21
12	Chi-Square Tests	22
13	Inference for Slopes and Tables	23

using *The Practice of Statistics for the AP Exam: 6th Edition* by Starnes and Tabor

1 Data Analysis

1.1 What is Statistics?

Definition 1 (Statistics). The science of collecting, analyzing, and drawing conclusions from data.

Data is collected from *individuals* about certain *variables*.

Definition 2 (Individual, Variable). **Individuals** are objects described in a dataset. Typically people, but not always.

Variables are attributes that can take different values for different individuals.

For example, *individuals* may be households, and *variables* may be region, number of people, household income, etc. It's important to distinguish between *categorical* and *quantitative* variables:

Definition 3 (Categorical and Quantitative Variables). **Categorical Variables** are variables whose values can be placed into distinct categories.

Quantitative Variables are variables whose values are quantities, typically counts or measurements.

For example, region would be categorical, while household income would be quantitative. *Not all numbers are quantitative*; eg. zip code.

1.2 Analyzing Categorical Data

1.2.1 One-Variable Categorical Data

Definition 4 (Frequency and Relative Frequency Tables). **Frequency Tables** shows the number of individuals that have values of a certain category. **Relative Frequency Tables** shows the proportion or percent of individuals in each category.

Note (relative) frequencies are not data; they summarize data. Bar graphs and Pie Graphs summarize relative frequency tables.

Beware of misleading graphs; we mainly react to the area of each bar, not the actual height.

	Like Skateboards	Do Not Like Skateboards	Totals
Like Snowmobiles	80	25	105
Do not like Snowmobiles	45	10	55
Totals	125	35	160

MathBits.com

Figure 1: An example two-way table with additional summary information.

1.2.2 Two-Variable Categorical Data

Use a two-way table to summarize data about two categorical variables. These tables can be used to answer questions about *marginal, joint, and conditional relative frequencies*.

Marginal relative frequencies give the percent or proportion of individuals that have a given value for one categorical variable. For example, the marginal relative frequency of liking skateboards is $\frac{125}{160} \approx 78.125\%$.

Joint relative frequencies give the percent or proportion of individuals that have a specific value for both categorical variables. For example, the joint frequency of liking both skateboards and snowmobiles is $\frac{80}{160} = 50\%$.

Conditional relative frequencies give the percent or proportion of individuals that have a specific value for one categorical variable relative to other individuals with the same other categorical variable. For example, the conditional relative frequency of those who like snowmobiles out of all individuals that like skateboards is $\frac{80}{125} = 64\%$.

These frequencies can be summarized in *side by side bar graphs, segmented bar graphs, or mosaic plots*.

Graphs and these tables can be used to show **association** between two variables. There is association between two variables if knowing the value of one helps to predict the other. For example, knowing that an individual likes skateboards helps predict whether they like snowmobiles ($\frac{80}{125} = 64\%$ vs $\frac{25}{35} \approx 71.4\%$). **ASSOCIATION DOES NOT IMPLY CAUSATION!**




Figure 2: A dotplot showing the distribution of kilometers run by members of the running club.

1.3 Analyzing Quantative Data with Graphs

Dotplots (as shown above) show each individual as a dot above their quantitative data value.

When describing the shape of a dotplot (or other quantitative graphs), *focus on main features*: major peaks, clusters, or gaps. Especially note whether the distribution is roughly symmetric or skewed:

Definition 5 (Symmetric, Skewed). A distribution is roughly **symmetric** if the right side of the graph has roughly the same shape as the left side.

A distribution is **skewed to the right** if the right 'tail' has less values than the left; typically, the left has a peak whereas the right does not. **Left-skewed** definition are defined similarly to right-skewed distributions.

For example, the distribution of the number of kilometers run is right-skewed because the right 'tail' has less values.

Graphs with a single peak are considered *unimodal*, like the dotplot. Distributions with two peaks are considered *bimodal*, and beyond that is considered *multimodal*.

When describing a distribution of quantitative data, use the acronym ROCS: **R**ange (max - min), **O**utliers (clear departures from the data), **C**enter (mean or median), and **S**hape (symmetry, skew, gaps, peaks).

Leaf plots exist. Stem represents first few digits, leaf represents final digit.




Figure 3: An example histogram with a normal distribution.

Histograms are a notable way of displaying quantative data, as they avoid showing individual data points. Histograms divide the variable into many 'bins' (bars), with the height representing the frequency. Smaller bins show more detail at the cost of a less clear pattern.

Don't confuse histograms and bar graphs. Histograms are used for quantative data, while bar graphs are used for qualatative data.

Use percentages when comapring to distributions in order to remove the effect of a larger sample.

1.4 Describing Quantative Data with Numbers

Definition 6 (Mean: \bar{x}, μ). The average of all individual data values. If there are n observations x_1, x_2, \dots, x_n , the sample mean is calculated by

$$\bar{x} = \frac{\sum x_i}{n}$$

The mean of a **sample** is referred to using \bar{x} , while the mean of a **population** is referred to using μ .

Statistics come from **samples** (small subset of population) and **parameters** come from **populations** (all possible samples of what's being tested).

The mean is not **resistant** as it is sensitive to strong outliers in a distribution. The median *is* a resistant measure of the center of the distribution.

Definition 7 (Median). The 'midpoint' of a distribution. Either the middle element (n is odd) or the average of the two middle elements (n is even) in a **SORTED** distribution.

Using both the mean and the median, one can predict the skew of the data. If a distribution is roughly symmetric without outliers, **the mean and median will be similar**. If a distribution is strongly skewed, **the mean will be pulled in the direction of the skew**. (Mean < Median for left-skewed, Mean > Median for right-skewed)

The **range** (max - min) is one way to show the variability of a distribution. Note that the range is *not* resistant.

Definition 8 (Standard Deviation). The **standard deviation** (s_x, σ) measures the 'average' distance of the values in a distribution from the mean. Standard deviation is calculated by

$$s_x = \sqrt{\frac{\sum (x_i - \bar{x})^2}{n - 1}}$$

The squared stdev is known as **variance** (s_x^2, σ^2). Remember, s_x refers to a sample while σ refers to a population. Larger stdev indicates greater variation, but is not a resistant measure of variability. **Stdev measures variance around the mean; if the mean is skewed, so will stdev!**

The **Interquartile Range (IQR)** is another way to measure variance, using $IQR = Q_3 - Q_1$ where Q represents the quartiles. IQR can be thought of as the range of the 'middle half' of the distribution. *IQR is a resistant measure.*

Lower Outliers < $Q_1 - 1.5 \times IQR$ or High Outliers > $Q_3 + 1.5 \times IQR$

boxplots and the five-number summary = min, Q1, median, Q3, and max exist. **Boxplots don't show gaps, clusters, or multiple peaks.**

be careful with language- 'skews' is a shape, IQR and range are single numbers (no 'in the middle of the IQR')

2 Modelling Distributions

3 Two-Variable Data

4 Collecting Data

5 Probability

6 Random Variables and Distributions

7 Sampling Distributions

8 Confidence Intervals

9 Significance Tests

Significance Tests are like the opposite of confidence intervals. Instead of using a statistic to find a parameter, significance tests use statistics to test claims about a parameter.

Definition 9 (Significance Test). A formal procedure for using observed statistics in order to decide between two competing claims (*hypotheses*) about parameters.

9.1 Basics of Significance Tests

9.1.1 Hypotheses

Definition 10 (Null Hypothesis, H_0). A claim about a parameter that we weigh evidence **against** in a significance test. Usually a statement of 'no difference' (as claimed).

Definition 11 (Alternative Hypothesis, H_a). The claim that we are trying to find evidence for. Directly contradicts the null hypothesis.

This "Null Hypothesis" and "Alternative Hypothesis" can be thought as trying to prove someone "guilty" or "not guilty."

For example, if a player claims they're a 80% free throw player, the null hypothesis would be $H_0 : p = 0.80$, and the alternative hypothesis would be $H_a : p < 0.80$.

The alternative hypothesis is **one-sided** because we suspect that the player makes less than 80% of his free throws. If we believe that it's equally plausible that they make more than 80% of their free throws, then we would use a **two-sided hypothesis**- $H_a : p \neq 0.80$.

Hypotheses express our beliefs before looking at the data. Molding a hypothesis around data shows nothing.

9.1.2 P-values

Definition 12 (P-value). The probability of getting the values observed in the data under the assumption that the null hypothesis H_0 is true.

Small P-values provide convincing evidence for the alternative hypothesis because small values suggest the observed result is unlikely to happen when the null hypothesis is true. Similarly, large P-values provide convincing evidence for the null hypothesis because large values suggest the observed result is likely to happen due to chance if the null hypothesis is true.

In terms of probability notation: P-value = $P(\text{observed data} \mid \text{null hypothesis is true})$.

For two-sided tests, we look at the distance between the null hypothesis and the observed data. For example, if $H_0 : p = 0.5$, $H_a : p \neq 0.5$ and $\hat{p} = 0.65$ (observed proportion), then the P-value = $P(\hat{p} \leq 0.35 \text{ or } \hat{p} \geq 0.6 \mid p = 0.5)$. We look at $\hat{p} \leq 0.35$ because $|p - 0.35| = |p - 0.65|$.

Based on the P-value, we make a conclusion about data:

- If the P-value is small (unlikely to happen by chance), then we "reject H_0 " and conclude that there is convincing evidence for H_a (in context).
- If the P-value is large (likely to happen by chance), then we "fail to reject H_0 " and conclude that there is not convincing evidence for H_a (in context).

How small does a P-value have to be in order to reject H_0 ? We use a given **significance value** for this boundary.

Definition 13 (Significance Level, α). The value that we use as a boundary for deciding whether a P-value is significant enough to disqualify the null hypothesis. α should be stated before data is produced (*cherrypicking*).

$\text{P-value} < \alpha \Rightarrow \text{reject } H_0 \Rightarrow \text{convincing evidence for } H_a \text{ in context}$

$\text{P-value} > \alpha \Rightarrow \text{fail to reject } H_0 \Rightarrow \text{not convincing evidence for } H_a \text{ in context}$

If P is less than the significance level, we say that the result is "statistically significant at the $\alpha = \text{---}$ level." Alternatively, "the results were significant ($P = 0.03 < \alpha = 0.05$)."
Keep in mind a P-value is more informative than a statement of significance!

NEVER "accept H_0 " or conclude that H_0 is true! Always use 'reject' or 'fail to reject!'

9.1.3 Type I and Type II Errors

When drawing a conclusion from a significance test, our conclusion may be wrong. There are two types we can make with the conclusion process, helpfully named "Type I" and "Type II" errors. Only one type of error is possible at once.

Definition 14 (Type I and II errors). A **Type I error** occurs when we reject H_0 when H_0 is true; the data gives convincing evidence that H_a is true despite being false.

A **Type II error** occurs when we fail to reject H_0 when H_a is true; the data fails to give convincing evidence that H_a is true despite it being true.

Note $P(\text{Type I error}) = \alpha$. However, significance is inversely proportional to Type II error: as significance decreases, $P(\text{Type II error})$ increases.

9.1.4 Power

Definition 15 (Power). The **power** of a test is the probability that the test finds convincing evidence for H_a given that the parameter is a specific value that does follow H_a . In other words:

$$\text{power} = P(\text{reject } H_0 \mid H_0 \text{ is false})$$

Power depends on a specific value. For example, if the power of a test to detect $p = 0.08$ is 0.29 given $H_0 : p = 0.10$, if the true proportion in the population is $p = 0.08$, there is a 0.29 probability that researchers will find convincing evidence for H_a .

Note $\text{Power} = 1 - P(\text{Type II Error})$, or written alternatively, $P(\text{Type II Error}) = 1 - \text{Power}$. Also note that the power of a test increases with higher sample number, higher significance level, further null and alternative parameter values, and 'wise choices when collecting data' (reducing variability). *Larger significance values reduce the chance for Type II error, but increase the chance for Type I error.*

9.1.5 Steps for Significance Tests

Overall, significance tests can be summarized in four steps:

1. **STATE** the hypotheses, significance level, and parameters.
2. **PLAN** the appropriate inference method and check conditions.
3. **DO** the significance test (see Section 7.2):
 - Give the sample statistic(s)
 - Calculate the standardized test statistic
 - Find the P-value.
4. **CONCLUDE** whether or not you believe in the null or alternate hypotheses, with reference to P-values, in the context of the problem. Make sure to reference the parameter, not the sample statistic!

9.2 Tests About a Population Proportion

Like confidence intervals, significance tests should satisfy several conditions; random sampling, no bias, 10% if applicable ($n < 0.10N$), and large counts ($np_0 \geq 10$ and $n(1 - p_0) \geq 10$). Note that the large counts condition uses p_0 : the number of successes and failures *assuming the null hypothesis is true* are both greater than or equal to 10.

To conduct a **one-sample z test for a proportion** (CITE THIS WHEN FOLLOWING THE SIGNIFIANCE TEST STEPS) about a proportion, look at the normal distribution *assuming the null hypothesis is true*.

From Section 7.2 (add ref later), $\mu_{\hat{p}} = p_0$ and $\sigma_{\hat{p}} = \sqrt{\frac{p_0(1-p_0)}{n}}$.

Using this, we standardize the statistic with respect to the null value to get the **standardized test statistic**; how many stdev units away the sample statistic is from what we would expect, assuming the null hypothesis is true. **YOU HAVE TO PUT THE STANDARDIZED TEST STATISTIC ON THE TEST!**

$$\text{standardized test statistic} = z_{sts} = \frac{\hat{p} - \mu_{\hat{p}}}{\sigma_{\hat{p}}}$$

Using this, we can directly find the P-value using `normalcdf` with the appropriate values:

- $P(z > z_{sts})$ if $H_a : p > p_0$
- $P(z < z_{sts})$ if $H_a : p < p_0$
- $P(z > |z_{sts}|) + P(z < -|z_{sts}|)$ if $H_a : p \neq p_0$

Keep in mind that just because the P-value is low *does not prove that H_0 is false*; sample proportions may be small due to sampling variability, or we have made a Type I error somehow. **AVOID JUMPING TO CONCLUSIONS**: just because a sample statistic looks unlikely, doesn't mean its P-value is!

It should be noted there is a link between *two-sided tests* and confidence intervals: a K% confidence interval gives an approximate set of p_0 's that would not be rejected by a two-sided test at the $\alpha = \frac{K}{100}$ level, with all other values being rejected.

9.3 Tests About a Difference in Proportions

Tests about difference in proportions work very similarly to tests about single proportions. For null hypothesis $H_0 : p_1 - p_2 = 0$, large counts can't directly

be applied; we need estimate the true difference in proportions p using a weighted average of the two proportions:

$$\hat{p}_C = \frac{\text{number of successes in both samples combined}}{\text{number of individuals in both samples combined}} = \frac{n_1\hat{p}_1 + n_2\hat{p}_2}{n_1 + n_2}$$

In this case, we use the combined proportion for large counts *for both samples*: that is, $n_1\hat{p}_C, n_1(1 - \hat{p}_C), n_2\hat{p}_C, n_2(1 - \hat{p}_C)$ all need to be greater or equal than 10. Using the combined proportion works for a non-zero difference in the null hypothesis, although is not strictly necessary (large counts for each individual distribution can be used).

The standardized test statistic also uses the combined estimate (for both proportions are equal; otherwise, each individual statistic can be used):

$$z = \frac{(\hat{p}_1 - \hat{p}_2) - 0}{\sqrt{\frac{\hat{p}_C(1-\hat{p}_C)}{n_1} + \frac{\hat{p}_C(1-\hat{p}_C)}{n_2}}}$$

10 Estimating Means

11 Confidence with Means

12 Chi-Square Tests

13 Inference for Slopes and Tables

$$\begin{aligned}\prod_{k=1}^4(x_k^2 + 1) &= \prod_{k=1}^4(x_k - i)(x_k + i) = \prod_{k=1}^4(x_k + i) \times \prod_{k=1}^4(x_k - i) \\ &= P(i)P(-i) = (b - d + 1)^2 + (a - c)^2 = (5 - 1)^2 + 0 = \boxed{16}\end{aligned}$$